

```
In [280]: import pandas as pd
import numpy as np
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
import matplotlib.pyplot as plt
```

```
In [281]: train= pd.read_csv(r"C:\Users\Ravi\Downloads\titanic\train.csv")
```

```
In [282]: train.head()
```

```
Out[282]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	Na
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C8
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	Na
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C12
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	Na

In [283]: `train.describe(include="all")`

Out[283]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch
<b>count</b>	891.000000	891.000000	891.000000	891	891	714.000000	891.000000	891.000000
<b>unique</b>	NaN	NaN	NaN	891	2	NaN	NaN	NaN
<b>top</b>	NaN	NaN	NaN	Andersson, Miss. Ingeborg Constanza	male	NaN	NaN	NaN
<b>freq</b>	NaN	NaN	NaN	1	577	NaN	NaN	NaN
<b>mean</b>	446.000000	0.383838	2.308642	NaN	NaN	29.699118	0.523008	0.381594
<b>std</b>	257.353842	0.486592	0.836071	NaN	NaN	14.526497	1.102743	0.806057
<b>min</b>	1.000000	0.000000	1.000000	NaN	NaN	0.420000	0.000000	0.000000
<b>25%</b>	223.500000	0.000000	2.000000	NaN	NaN	20.125000	0.000000	0.000000
<b>50%</b>	446.000000	0.000000	3.000000	NaN	NaN	28.000000	0.000000	0.000000
<b>75%</b>	668.500000	1.000000	3.000000	NaN	NaN	38.000000	1.000000	0.000000
<b>max</b>	891.000000	1.000000	3.000000	NaN	NaN	80.000000	8.000000	6.000000

In [284]: `train.columns`

Out[284]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'], dtype='object')

In [285]: `x= train.drop("Survived",axis=1)`

In [286]: `y= train["Survived"]`

In [287]: `from sklearn.model_selection import train_test_split`

In [288]: `train_x, test_x, train_y, test_y = train_test_split(x,y,random_state = 40, stratify=y)`

In [289]: `train_x.shape`

Out[289]: (668, 11)

In [290]: `test_x.shape`

Out[290]: (223, 11)


## Univariate Analysis

```
In [291]: train_x["Sex"].value_counts()
```

```
Out[291]: male      432  
female    236  
Name: Sex, dtype: int64
```


```
In [292]: train_x["Sex"].value_counts().plot(kind="bar")
```

```
Out[292]: <matplotlib.axes._subplots.AxesSubplot at 0x16135893b08>
```




```
In [293]: train_x["Age"].plot.box()
```

```
Out[293]: <matplotlib.axes._subplots.AxesSubplot at 0x161381f3dc8>
```




```
In [294]: sns.distplot(train_x["Age"], hist = True, kde=True)
```

```
Out[294]: <matplotlib.axes._subplots.AxesSubplot at 0x16138258a48>
```



```
In [295]: train_x["SibSp"].value_counts().plot.bar()
```

```
Out[295]: <matplotlib.axes._subplots.AxesSubplot at 0x161382e3d08>
```




```
In [296]: train_x["Pclass"].value_counts()
```

```
Out[296]: 3    376  
          1    158  
          2    134  
          Name: Pclass, dtype: int64
```


```
In [297]: train_x["Pclass"].value_counts().plot.bar()
```

```
Out[297]: <matplotlib.axes._subplots.AxesSubplot at 0x16138363448>
```



```
In [298]: train_x["Fare"].plot.hist(color="Green")
```

```
Out[298]: <matplotlib.axes._subplots.AxesSubplot at 0x161383bfd48>
```



```
In [299]: train_x["Cabin"].value_counts().sum()
```


```
Out[299]: 147
```

```
In [300]: train_x["Cabin"].isna().value_counts()
```

```
Out[300]: True      521  
         False    147  
         Name: Cabin, dtype: int64
```

```
In [301]: train_x["Embarked"].value_counts().plot.hist(color= "Magenta")
```

```
Out[301]: <matplotlib.axes._subplots.AxesSubplot at 0x16138431348>
```



## Checking for missing values

```
In [302]: train_x.isnull().sum()
```

```
Out[302]: PassengerId      0
Pclass      0
Name        0
Sex         0
Age        136
SibSp       0
Parch       0
Ticket      0
Fare        0
Cabin      521
Embarked     1
dtype: int64
```

Therefore values from "Age" and "Cabin" are missing, which are needed to be taken care of.

```
In [303]: mn= train["Age"].mean()
```

```
In [304]: train["Age"]=train_x["Age"].replace(np.nan,mn,inplace = True)
```

Here new methods has to checked after building the model.

1. Dropping the rows
2. Creating a new dimension


```
In [305]: train_x.head()
```

```
Out[305]:
```

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
660	661	1	Frauenthal, Dr. Henry William	male	50.000000	2	0	PC 17611	133.6500	NaN
576	577	2	Garside, Miss. Ethel	female	34.000000	0	0	243880	13.0000	NaN
158	159	3	Smiljanic, Mr. Mile	male	29.699118	0	0	315037	8.6625	NaN
617	618	3	Lobb, Mrs. William Arthur (Cordelia K Stanlick)	female	26.000000	1	0	A/5. 3336	16.1000	NaN
407	408	2	Richards, Master. William Rowe	male	3.000000	1	1	29106	18.7500	NaN

```
In [306]: train_x["Age"].plot.box()
```

```
Out[306]: <matplotlib.axes._subplots.AxesSubplot at 0x1613846c188>
```



## Bivariate Analysis

```
In [307]: train_x.shape
```

```
Out[307]: (668, 11)
```

```
In [308]: test_x.shape
```

```
Out[308]: (223, 11)
```


```
In [309]: train.columns
```

```
Out[309]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
                'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
               dtype='object')
```

```
In [310]: sex_sur=pd.crosstab(train_x["Sex"],train_y)
```

```
In [311]: sex_sur.div(sex_sur.sum(1).astype(float),axis=0).plot.bar(stacked=True, figsize=(
```

```
Out[311]: <matplotlib.axes._subplots.AxesSubplot at 0x161384da2c8>
```




Chances of female survival are more than the males

```
In [312]: p_sur=pd.crosstab(train_x["Pclass"],train_y)
```




```
In [313]: p_sur.plot.bar(stacked=True,figsize=(6,6))
```

```
Out[313]: <matplotlib.axes._subplots.AxesSubplot at 0x1613854e208>
```



```
In [314]: age_fare=pd.crosstab(train_x["Age"],train_x["Fare"])
```

```
In [315]: train.plot.scatter("Age","Fare")  
plt.show()
```



```
In [316]: train_x["Age"].dtype
```

```
Out[316]: dtype('float64')
```

```
In [317]: #age Fare relation  
train["Age"].corr(train_x["Fare"])
```

```
Out[317]: nan
```

Therefore Age and Fare are not much related.

```
In [318]: #Age and survival  
train_x["Age"].corr(train_y)
```

```
Out[318]: -0.10188494983187488
```

```
In [319]: #fare and survival  
train_x["Fare"].corr(train_y)
```

```
Out[319]: 0.25569898233277427
```

```
In [320]: train_x["Pclass"].corr(train_y)
```

```
Out[320]: -0.3386341552390735
```

```
In [321]: corr= train.corr()  
fig,ax=plt.subplots()  
fig.set_size_inches(10,10)  
sns.heatmap(corr,annot=True,cmap="YlGnBu")
```

```
Out[321]: <matplotlib.axes._subplots.AxesSubplot at 0x1613863eec8>
```



## Model Building

### Logistic Regression

In [ ]:

```
In [322]: trainn=train.drop(["Survived"],axis=1)
trainn=train.drop(["Cabin"],axis=1)
target=train["Survived"]
```


```
In [323]: trainn=pd.get_dummies(trainn)
```

```
In [324]: trainn.head()
```

Out[324]:

	PassengerId	Survived	Pclass	SibSp	Parch	Fare	Name_Abbing, Mr. Anthony	Name_Abbott, Mr. Rossmore Edward	Name_Ab Mrs. Sta (Rosa H
0	1	0	3	1	0	7.2500	0	0	
1	2	1	1	1	0	71.2833	0	0	
2	3	1	3	0	0	7.9250	0	0	
3	4	1	1	1	0	53.1000	0	0	
4	5	0	3	0	0	8.0500	0	0	

5 rows × 1583 columns



```
In [325]: train.columns
```

```
Out[325]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
                'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
                dtype='object')
```

```
In [326]: trainn=trainn.drop(["Survived"],axis=1)
```

In [327]: `trainn.head()`

Out[327]:

	PassengerId	Pclass	SibSp	Parch	Fare	Name_Abbing, Mr. Anthony	Name_Abbott, Mr. Rossmore Edward	Name_Abbott, Mrs. Stanton (Rosa Hunt)	Nam
0	1	3	1	0	7.2500	0	0	0	
1	2	1	1	0	71.2833	0	0	0	
2	3	3	0	0	7.9250	0	0	0	
3	4	1	1	0	53.1000	0	0	0	
4	5	3	0	0	8.0500	0	0	0	

5 rows × 1582 columns

In [328]: `target.head()`

Out[328]:

```
0    0
1    1
2    1
3    1
4    0
Name: Survived, dtype: int64
```

In [329]: `train_X, test_X, train_Y, test_Y = train_test_split(trainn, target, test_size=0.3, rand`

In [330]: `from sklearn.linear_model import LogisticRegression`

In [331]: `lrg=LogisticRegression()`

In [332]: `test_X.head()`

Out[332]:

	PassengerId	Pclass	SibSp	Parch	Fare	Name_Abbing, Mr. Anthony	Name_Abbott, Mr. Rossmore Edward	Name_Abbott, Mrs. Stanton (Rosa Hunt)	Ni
646	647	3	0	0	7.8958	0	0	0	
876	877	3	0	0	9.8458	0	0	0	
359	360	3	0	0	7.8792	0	0	0	
194	195	1	0	0	27.7208	0	0	0	
819	820	3	3	2	27.9000	0	0	0	

5 rows × 1582 columns

```
In [333]: lrg.fit(train_X,train_Y)
```

```
Out[333]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='auto', n_jobs=None, penalty='l2',
                             random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                             warm_start=False)
```

```
In [334]: pred=lrg.predict(test_X)
```

```
In [335]: pred.shape
```

```
Out[335]: (268,)
```

```
In [336]: from sklearn.metrics import accuracy_score
```

```
In [337]: accuracy_score(test_Y,pred)
```

```
Out[337]: 0.8134328358208955
```

## Decision Tree

```
In [338]: from sklearn.tree import DecisionTreeClassifier
```

```
In [339]: clf=DecisionTreeClassifier(max_depth=16,random_state=40)
```

```
In [340]: clf.fit(train_X,train_Y)
```

```
Out[340]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                                 max_depth=16, max_features=None, max_leaf_nodes=None,
                                 min_impurity_decrease=0.0, min_impurity_split=None,
                                 min_samples_leaf=1, min_samples_split=2,
                                 min_weight_fraction_leaf=0.0, presort='deprecated',
                                 random_state=40, splitter='best')
```

```
In [341]: pred1=clf.predict(test_X)
```

```
In [342]: accuracy_score(test_Y,pred1)
```

```
Out[342]: 0.8507462686567164
```

**After tuning the parameters a an accuracy of 85.07% is achieved.**

```
In [343]: #predicting on training data set
pre=clf.predict(train_X)
```

```
In [344]: accuracy_score(train_Y,pre)
```

```
Out[344]: 0.9775280898876404
```

**On predicting on the dataset an accuracy of 97.75% is achieved.**

```
In [ ]:
```

### Working with Test Data Set

```
In [345]: test_set=pd.read_csv(r"C:\Users\Ravi\Downloads\titanic\test.csv")
```

```
In [346]: test_set.head()
```

```
Out[346]:
```

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	892	3	Kelly, Mr. James	male	34.5	0	0	330911	7.8292	NaN	C
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	363272	7.0000	NaN	S
2	894	2	Myles, Mr. Thomas Francis	male	62.0	0	0	240276	9.6875	NaN	C
3	895	3	Wirz, Mr. Albert	male	27.0	0	0	315154	8.6625	NaN	S
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1	3101298	12.2875	NaN	S

```
In [347]: test_set["Age"].replace(np.nan,mn,inplace=True)
```

```
In [348]: test_set.drop(["Cabin"],axis=1,inplace= True)
```

```
In [349]: test1=pd.get_dummies(test_set)
```

```
In [350]: test_set[test_set["Fare"].isna()]= 8
```

```
In [351]: test_set.isna().sum()
```

```
Out[351]: PassengerId    0
          Pclass        0
          Name          0
          Sex           0
          Age           0
          SibSp         0
          Parch         0
          Ticket        0
          Fare          0
          Embarked      0
          dtype: int64
```

```
In [ ]:
```

```
In [352]: train_X.shape
```

```
Out[352]: (623, 1582)
```

```
In [353]: test1.shape
```

```
Out[353]: (418, 792)
```

```
In [354]: train.columns
```

```
Out[354]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
                 'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
                 dtype='object')
```

```
In [355]: test_set.columns
```

```
Out[355]: Index(['PassengerId', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch',
                 'Ticket', 'Fare', 'Embarked'],
                 dtype='object')
```

## New Attempt



In [356]: `train.head()`

Out[356]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cat
0	1	0	3	Braund, Mr. Owen Harris	male	None	1	0	A/5 21171	7.2500	Na
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	None	1	0	PC 17599	71.2833	C
2	3	1	3	Heikkinen, Miss. Laina	female	None	0	0	STON/O2. 3101282	7.9250	Na
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	None	1	0	113803	53.1000	C1
4	5	0	3	Allen, Mr. William Henry	male	None	0	0	373450	8.0500	Na

In [ ]:

In [357]: `features=["Pclass", "Sex", "Age", "SibSp", "Parch", "Fare", "Embarked"]`

In [358]: `df=train[features]`

In [359]: `df.head()`

Out[359]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	male	None	1	0	7.2500	S
1	1	female	None	1	0	71.2833	C
2	3	female	None	0	0	7.9250	S
3	1	female	None	1	0	53.1000	S
4	3	male	None	0	0	8.0500	S

In [360]: `df["Age"].replace(np.nan, mn, inplace=True)`

In [361]: `df.shape`

Out[361]: (891, 7)

```
In [362]: df1=pd.get_dummies(df)
```

```
In [363]: df1.head()
```

```
Out[363]:
```

	Pclass	Age	SibSp	Parch	Fare	Sex_female	Sex_male	Embarked_C	Embarked_Q	E
0	3	29.699118	1	0	7.2500	0	1	0	0	
1	1	29.699118	1	0	71.2833	1	0	1	0	
2	3	29.699118	0	0	7.9250	1	0	0	0	
3	1	29.699118	1	0	53.1000	1	0	0	0	
4	3	29.699118	0	0	8.0500	0	1	0	0	

```
In [364]: df1.shape
```

```
Out[364]: (891, 10)
```

```
In [365]: trn_x,tst_x,trn_y,tst_y=train_test_split(df1,target,test_size=0.3)
```

```
In [366]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import PolynomialFeatures
```

```
In [447]: Input=[("scale",StandardScaler()),("polynomial",PolynomialFeatures(degree=2,inclu
```

```
In [448]: pipe=Pipeline(Input)
```

```
In [449]: pipe.fit(trn_x,trn_y)
```

```
Out[449]: Pipeline(memory=None,
                 steps=[('scale',
                        StandardScaler(copy=True, with_mean=True, with_std=True)),
                        ('polynomial',
                        PolynomialFeatures(degree=2, include_bias=False,
                                           interaction_only=False, order='C')),
                        ('mode',
                        LogisticRegression(C=1.0, class_weight=None, dual=False,
                                           fit_intercept=True, intercept_scaling=1,
                                           l1_ratio=None, max_iter=100,
                                           multi_class='auto', n_jobs=1, penalty='l2',
                                           random_state=10, solver='lbfgs', tol=0.000
1,
                                           verbose=0, warm_start=False))],
                 verbose=False)
```

```
In [450]: pred_val=pipe.predict(tst_x)
```